


## Assessment 2

**Problem 2: Try different hyperparameters and perform convolution transform on an image**

### Convolution operation on an image

The convolution layer (CONV) uses filters that perform convolution operations as it is scanning the input  $I$  with respect to its dimensions. Its hyperparameters include the filter size  $F$  and stride  $S$ . The resulting output  $O$  is called feature map or activation map.



In [2]:

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
import skimage.measure
```

**Loading the image form local directory and checking the entropy of the image**


In [3]:


```
input_image = cv2.imread('/content/drive/MyDrive/DumpYard/paris-843229_1920.jpg')
input_image = cv2.cvtColor(input_image, cv2.COLOR_BGR2RGB)
fig, ax = plt.subplots(1, figsize=(9,6))
entropy1 = skimage.measure.shannon_entropy(input_image)
print("Entropy for the original image is:",entropy1)
plt.imshow(input_image)
```

Entropy for the original image is: 7.4473113904233825

Out[3]:

<matplotlib.image.AxesImage at 0x7f257706e150>





## Use a simple 3x3 filter on the image and demonstrate the transformed image with its entropy

### Filter

A filter of size  $F \times F$  applied to an input containing  $C$  channels is a  $F \times F \times C$  volume that performs convolutions on an input of size  $I \times I \times C$  and produces an output feature map (also called activation map) of size  $O \times O \times 1$ .

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

\*=

### Kernel

In [4]:

```
sharpenfilter = np.array([[0, -1, 0], [-1, 9, -1], [0, -1, 0]], np.float32)/9
sharpenfilter
```

Out[4]:

```
array([[ 0., -0.11111111,  0.],
       [-0.11111111,  1., -0.11111111],
       [ 0., -0.11111111,  0.]], dtype=float32)
```

### Transformation and output image with its Entropy

In [5]:

```
sharpen = cv2.filter2D(src=input_image, kernel=sharpenfilter, ddepth=-1)
fig, ax = plt.subplots(1, figsize=(9,6))
entropy2 = skimage.measure.shannon_entropy(sharpen)
print("Entropy for the sharpened image is:",entropy2)
plt.imshow(sharpen)
```

Entropy for the sharpened image is: 6.647542774487621

Entropy for the sharpened image is: 0.017512771107021

Out[5]:

<matplotlib.image.AxesImage at 0x7f25762e1190>



**Use a simple 5x5 filter on the image and demonstrate the transformed image with its entropy**

#### Filter or Kernel

In [6]:

```
gaussian_blur_filter = np.array([[[1,4,6,4,1], [4,16,24,16,4], [6,24,36,24,6], [4,16,24,16,4], [1,4,6,4,1]]], np.float32)/256
gaussian_blur_filter
```

Out[6]:

```
array([[0.00390625, 0.015625, 0.0234375, 0.015625, 0.00390625],
       [0.015625, 0.0625, 0.09375, 0.0625, 0.015625],
       [0.0234375, 0.09375, 0.140625, 0.09375, 0.0234375],
       [0.015625, 0.0625, 0.09375, 0.0625, 0.015625],
       [0.00390625, 0.015625, 0.0234375, 0.015625, 0.00390625]],
      dtype=float32)
```

#### Transformation and Output image with Entropy


In [7]:


```
gaussianBlur = cv2.filter2D(src=input_image, kernel=gaussian_blur_filter, ddepth=-1)
entropy3 = skimage.measure.shannon_entropy(gaussianBlur)
print("Entropy for the blurred image is:",entropy3)
fig, ax = plt.subplots(1, figsize=(9,6))
plt.imshow(gaussianBlur)
```

Entropy for the blurred image is: 7.395531620068165

Out[7]:

<matplotlib.image.AxesImage at 0x7f25746579d0>






In [8]:

```
grey_image = cv2.cvtColor(input_image, cv2.COLOR_BGR2GRAY)
```

## Filtering an image with 3x3 filter with a stride = 2 and then calculate its Entropy

### Stride

For a convolutional or a pooling operation, the stride  $S$  denotes the number of pixels by which the window moves after each operation.



In [9]:

```
def convolve2D(image, kernel, strides):
    kernel = np.flipud(np.fliplr(kernel))
    xKernShape = kernel.shape[0]
    yKernShape = kernel.shape[1]
    xImgShape = image.shape[0]
    yImgShape = image.shape[1]
    xOutput = int((xImgShape - xKernShape + 2) / strides) + 1
    yOutput = int((yImgShape - yKernShape + 2) / strides) + 1
    output = np.zeros((xOutput, yOutput))
    for y in range(image.shape[1]):
        if y > image.shape[1] - yKernShape:
            break
        if y % strides == 0:
            for x in range(image.shape[0]):
                if x > image.shape[0] - xKernShape:
                    break
                try:
                    if x % strides == 0:
                        output[x, y] = (kernel * image[x: x + xKernShape, y: y + yKernS
hape]).sum()
                except:
                    break
    return output


kernel = np.array([[0, -1, 0], [-1, 9, -1], [0, -1, 0]], np.float32)/9
optimg = convolve2D(grey_image, kernel, strides = 2)
entropy6 = skimage.measure.shannon_entropy(optimg)
print("Entropy for an image with kernel = 3x3 and stride = 2 is:", entropy6)
fig, ax = plt.subplots(1, figsize=(9,6))
plt.imshow(optimg)
```

Entropy for an image with kernel = 3x3 and stride = 2 is: 3.222862544247595

Out[9]:

<matplotlib.image.AxesImage at 0x7f25745e660>

<matplotlib.image.AxesImage at 0x7125745ac690>



## Filtering an image with 5x5 filter with a stride = 2 and then calculate its Entropy

In [10]:

```
def convolve2D(image, kernel, strides):
    kernel = np.flipud(np.fliplr(kernel))
    xKernShape = kernel.shape[0]
    yKernShape = kernel.shape[1]
    xImgShape = image.shape[0]
    yImgShape = image.shape[1]
    xOutput = int((xImgShape - xKernShape + 2) / strides) + 1
    yOutput = int((yImgShape - yKernShape + 2) / strides) + 1
    output = np.zeros((xOutput, yOutput))
    for y in range(image.shape[1]):
        if y > image.shape[1] - yKernShape:
            break
        if y % strides == 0:
            for x in range(image.shape[0]):
                if x > image.shape[0] - xKernShape:
                    break
                try:
                    if x % strides == 0:
                        output[x, y] = (kernel * image[x: x + xKernShape, y: y + yKernS
hape]).sum()
                except:
                    break
    return output


kernel = np.array([[1,4,6,4,1], [4,16,24,16,4], [6,24,36,24,6], [4,16,24,16,4], [1,4,6,4,1]]), np.float32)/256
optimg = convolve2D(grey_image, kernel, strides = 2)
entropy7 = skimage.measure.shannon_entropy(optimg)
print("Entropy for an image with kernel = 5x5 and stride = 2 is:",entropy7)
fig, ax = plt.subplots(1, figsize=(9,6))
plt.imshow(optimg)
```

Entropy for an image with kernel = 5x5 and stride = 2 is: 4.189303058746414

Out[10]:

<matplotlib.image.AxesImage at 0x7f2574581b90>





## Filtering an image with 3x3 filter with stride = 1 and zero padding, then calculate its Entropy

### Zero Padding

Zero-padding denotes the process of adding  $P$  zeroes to each side of the boundaries of the input.

0	0	0	0	0	0
0	35	19	25	6	0
0	13	22	16	53	0
0	4	3	7	10	0
0	9	8	1	3	0
0	0	0	0	0	0

In [11]:


```
def convolve2d(image, kernel):
    output = np.zeros_like(image)
    image_padded = np.zeros((image.shape[0] + 2, image.shape[1] + 2))
    image_padded[1:-1, 1:-1] = image
    for x in range(image.shape[1]):
        for y in range(image.shape[0]):
            output[y, x] = (kernel * image_padded[y:y+3, x:x+3]).sum()
    return output
# For a 3x3 Kernel and stride = 1 and Zero_padding
KERNEL = np.array([[0, -1, 0], [-1, 9, -1], [0, -1, 0]], np.float32)/9
stride_1_3x3 = convolve2d(grey_image, kernel=KERNEL)
entropy4 = skimage.measure.shannon_entropy(stride_1_3x3)
print("Entropy for the 3x3 image with zero padding and stride = 1 is:", entropy4)
fig, ax = plt.subplots(1, figsize=(9,6))
plt.imshow(stride_1_3x3)
```

Entropy for the 3x3 image with zero padding and stride = 1 is: 6.055754583579611

Out[11]:

<matplotlib.image.AxesImage at 0x7f25745972d0>





## Filtering an image with 5x5 filter with stride = 1 and zero padding, then calculate its Entropy

In [12]:

```
def convolve2d(image, kernel):
    output = np.zeros_like(image)
    image_padded = np.zeros((image.shape[0] + 4, image.shape[1] + 4))
    image_padded[2:-2, 2:-2] = image
    for x in range(image.shape[1]):
        for y in range(image.shape[0]):
            output[y, x] = (kernel * image_padded[y:y+5, x:x+5]).sum()
    return output
# For a 5x5 Kernel and stride = 1 and Zero_padding
KERNEL = np.array([[1,4,6,4,1], [4,16,24,16,4], [6,24,36,24,6], [4,16,24,16,4], [1,4,6,4,1]]), np.float32)/256
opimg = convolve2d(grey_image, kernel=KERNEL)
entropy5 = skimage.measure.shannon_entropy(opimg)
print("Entropy for the 5x5 image with zero padding and stride = 1 is:", entropy5)
fig, ax = plt.subplots(1, figsize=(9,6))
plt.imshow(opimg)
```

Entropy for the 5x5 image with zero padding and stride = 1 is: 6.715918610760809

Out[12]:

<matplotlib.image.AxesImage at 0x7f257459bc50>

