

Assessment 6

GitHub Link: https://github.com/manansuthar55/CSE6037_20MAI0016/tree/main/Assessment_6

Image Colorization Using Autoencoders

In [1]:

```
import os
import sys
import random
import warnings
import numpy as np
import pandas as pd
import cv2
import matplotlib.pyplot as plt
from tqdm import tqdm
from itertools import chain
import skimage
from PIL import Image
from skimage.io import imread, imshow, imread_collection, concatenate_images
from skimage.transform import resize
from skimage.util import crop, pad
from skimage.morphology import label
from skimage.color import rgb2gray, gray2rgb, rgb2lab, lab2rgb
from sklearn.model_selection import train_test_split
from keras.applications.inception_resnet_v2 import InceptionResNetV2, preprocess_input
from keras.models import Model, load_model, Sequential
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Input, Dense, UpSampling2D, RepeatVector, Reshape
from keras.layers.core import Dropout, Lambda
from keras.layers.convolutional import Conv2D, Conv2DTranspose
from keras.layers.pooling import MaxPooling2D
from keras.layers.merge import concatenate
from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
from keras import backend as K
import tensorflow as tf

warnings.filterwarnings('ignore', category=UserWarning, module='skimage')
seed = 42
random.seed = seed
np.random.seed = seed
```

Using TensorFlow backend.

Read in Data

The dataset I am using is ~2000 classic paintings which we will remove the color from and attempt to teach a neural network to recolorize them.

In [2]:

```
IMG_WIDTH = 256
IMG_HEIGHT = 256
IMG_CHANNELS = 3
INPUT_SHAPE=(IMG_HEIGHT, IMG_WIDTH, 1)
TRAIN_PATH = '../input/art-images-drawings-painting-sculpture-engraving/dataset/dataset_updated/training_set/painting/'
train_ids = next(os.walk(TRAIN_PATH))[2]
```

In [3]:

```
%%time
```

```
X_train = np.zeros((len(train_ids)-86, IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS), dtype=np.uint8)
missing_count = 0
print('Getting train images ... ')
sys.stdout.flush()
for n, id_ in tqdm(enumerate(train_ids), total=len(train_ids)):
    path = TRAIN_PATH + id_ + ''
    try:
        img = imread(path)
        img = resize(img, (IMG_HEIGHT, IMG_WIDTH), mode='constant', preserve_range=True)
        X_train[n-missing_count] = img
    except:
        missing_count += 1
X_train = X_train.astype('float32') / 255.
print("Total missing: " + str(missing_count))
```

Getting train images ...


100%|██████████| 2128/2128 [00:39<00:00, 54.10it/s]

Total missing: 86
CPU times: user 44.8 s, sys: 25.1 s, total: 1min 9s
Wall time: 40.4 s

In [4]:

```
imshow(X_train[5])
plt.show()
```

/opt/conda/lib/python3.6/site-packages/skimage/io/_plugins/matplotlib_plugin.py:51: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
out_of_range_float = (np.issubdtype(image.dtype, np.float) and
/opt/conda/lib/python3.6/site-packages/matplotlib/axes/_base.py:1428: MatplotlibDeprecationWarning: The 'box-forced' keyword argument is deprecated since 2.2.
" since 2.2.", cbook.mplDeprecation)



In [5]:

```
X_train, X_test = train_test_split(X_train, test_size=20, random_state=seed)
```

Create the Model

In [6]:

```
inception = InceptionResNetV2(weights=None, include_top=True)
inception.load_weights('../input/inception-resnet-v2-weights/inception_resnet_v2_weights_tf_dim_ordering_tf_kernels.h5')
inception.graph = tf.get_default_graph()
```

In [7]:

```
def Colorize():
```

```

embed_input = Input(shape=(1000,))
#Encoder
encoder_input = Input(shape=(256, 256, 1,))
encoder_output = Conv2D(128, (3,3), activation='relu', padding='same',strides=1)(encoder_input)
encoder_output = MaxPooling2D((2, 2), padding='same')(encoder_output)
encoder_output = Conv2D(128, (4,4), activation='relu', padding='same')(encoder_output)
encoder_output = Conv2D(128, (3,3), activation='relu', padding='same',strides=1)(encoder_output)
encoder_output = MaxPooling2D((2, 2), padding='same')(encoder_output)
encoder_output = Conv2D(256, (4,4), activation='relu', padding='same')(encoder_output)
encoder_output = Conv2D(256, (3,3), activation='relu', padding='same',strides=1)(encoder_output)
encoder_output = MaxPooling2D((2, 2), padding='same')(encoder_output)
encoder_output = Conv2D(256, (4,4), activation='relu', padding='same')(encoder_output)
encoder_output = Conv2D(256, (3,3), activation='relu', padding='same')(encoder_output)
encoder_output = Conv2D(256, (3,3), activation='relu', padding='same')(encoder_output)
encoder_output = Conv2D(256, (3,3), activation='relu', padding='same')(encoder_output)
#Fusion
fusion_output = RepeatVector(32 * 32)(embed_input)
fusion_output = Reshape([32, 32, 1000])(fusion_output)
fusion_output = concatenate([encoder_output, fusion_output], axis=3)
fusion_output = Conv2D(256, (1, 1), activation='relu', padding='same')(fusion_output)
#Decoder
decoder_output = Conv2D(128, (3,3), activation='relu', padding='same')(fusion_output)
decoder_output = Conv2D(64, (3,3), activation='relu', padding='same')(decoder_output)
decoder_output = UpSampling2D((2, 2))(decoder_output)
decoder_output = Conv2D(128, (3,3), activation='relu', padding='same')(decoder_output)
decoder_output = UpSampling2D((2, 2))(decoder_output)
decoder_output = Conv2D(64, (4,4), activation='relu', padding='same')(decoder_output)
decoder_output = Conv2D(64, (3,3), activation='relu', padding='same')(decoder_output)
decoder_output = Conv2D(32, (2,2), activation='relu', padding='same')(decoder_output)
decoder_output = Conv2D(2, (3, 3), activation='tanh', padding='same')(decoder_output)
decoder_output = UpSampling2D((2, 2))(decoder_output)
return Model(inputs=[encoder_input, embed_input], outputs=decoder_output)

model = Colorize()
model.compile(optimizer='adam', loss='mean_squared_error')
model.summary()

```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_3 (InputLayer)	(None, 256, 256, 1)	0	
=====			
conv2d_204 (Conv2D)	(None, 256, 256, 128)	1280	input_3[0][0]
=====			
max_pooling2d_5 (MaxPooling2D)	(None, 128, 128, 128)	0	conv2d_204[0][0]
=====			
conv2d_205 (Conv2D)	(None, 128, 128, 128)	262272	max_pooling2d_5[0][0]

conv2d_206 (Conv2D)	(None, 128, 128, 128)	147584	conv2d_205[0][0]
max_pooling2d_6 (MaxPooling2D)	(None, 64, 64, 128)	0	conv2d_206[0][0]
conv2d_207 (Conv2D)	(None, 64, 64, 256)	524544	max_pooling2d_6[0][0]
conv2d_208 (Conv2D)	(None, 64, 64, 256)	590080	conv2d_207[0][0]
max_pooling2d_7 (MaxPooling2D)	(None, 32, 32, 256)	0	conv2d_208[0][0]
conv2d_209 (Conv2D)	(None, 32, 32, 256)	1048832	max_pooling2d_7[0][0]
input_2 (InputLayer)	(None, 1000)	0	
conv2d_210 (Conv2D)	(None, 32, 32, 256)	590080	conv2d_209[0][0]
repeat_vector_1 (RepeatVector)	(None, 1024, 1000)	0	input_2[0][0]
conv2d_211 (Conv2D)	(None, 32, 32, 256)	590080	conv2d_210[0][0]
reshape_1 (Reshape)	(None, 32, 32, 1000)	0	repeat_vector_1[0][0]
concatenate_1 (Concatenate)	(None, 32, 32, 1256)	0	conv2d_211[0][0] reshape_1[0][0]
conv2d_212 (Conv2D)	(None, 32, 32, 256)	321792	concatenate_1[0][0]
conv2d_213 (Conv2D)	(None, 32, 32, 128)	295040	conv2d_212[0][0]
conv2d_214 (Conv2D)	(None, 32, 32, 64)	73792	conv2d_213[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 64, 64, 64)	0	conv2d_214[0][0]
conv2d_215 (Conv2D)	(None, 64, 64, 128)	73856	up_sampling2d_1[0][0]


```

                                verbose=1,
                                factor=0.5,
                                min_lr=0.00001)

filepath = "Art_Colorization_Model.h5"
checkpoint = ModelCheckpoint(filepath,
                             save_best_only=True,
                             monitor='loss',
                             mode='min')

model_callbacks = [learning_rate_reduction,checkpoint]

```

Train the Model

In [10]:

```

%%time
BATCH_SIZE = 20
model.fit_generator(image_a_b_gen(X_train,BATCH_SIZE),
                    epochs=30,
                    verbose=1,
                    steps_per_epoch=X_train.shape[0]/BATCH_SIZE,
                    callbacks=model_callbacks
                    )

```

```

Epoch 1/30
102/101 [=====] - 174s 2s/step - loss: 0.0107
Epoch 2/30
102/101 [=====] - 159s 2s/step - loss: 0.0051
Epoch 3/30
102/101 [=====] - 164s 2s/step - loss: 0.0048
Epoch 4/30
102/101 [=====] - 164s 2s/step - loss: 0.0047
Epoch 5/30
102/101 [=====] - 164s 2s/step - loss: 0.0046
Epoch 6/30
102/101 [=====] - 166s 2s/step - loss: 0.0047
Epoch 7/30
102/101 [=====] - 163s 2s/step - loss: 0.0045
Epoch 8/30
102/101 [=====] - 164s 2s/step - loss: 0.0044
Epoch 9/30
102/101 [=====] - 164s 2s/step - loss: 0.0043
Epoch 10/30
102/101 [=====] - 163s 2s/step - loss: 0.0042
Epoch 11/30
102/101 [=====] - 162s 2s/step - loss: 0.0042
Epoch 12/30
102/101 [=====] - 162s 2s/step - loss: 0.0042
Epoch 13/30
102/101 [=====] - 161s 2s/step - loss: 0.0041
Epoch 14/30
102/101 [=====] - 163s 2s/step - loss: 0.0041
Epoch 15/30
102/101 [=====] - 162s 2s/step - loss: 0.0040
Epoch 16/30
102/101 [=====] - 162s 2s/step - loss: 0.0040
Epoch 17/30
102/101 [=====] - 162s 2s/step - loss: 0.0040
Epoch 18/30
102/101 [=====] - 162s 2s/step - loss: 0.0039
Epoch 19/30
102/101 [=====] - 162s 2s/step - loss: 0.0039
Epoch 20/30
102/101 [=====] - 162s 2s/step - loss: 0.0039
Epoch 21/30
102/101 [=====] - 161s 2s/step - loss: 0.0039

Epoch 00021: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
Epoch 22/30
102/101 [=====] - 161s 2s/step - loss: 0.0037
Epoch 23/30

```

```

102/101 [=====] - 162s 2s/step - loss: 0.0036
Epoch 24/30
102/101 [=====] - 162s 2s/step - loss: 0.0036
Epoch 25/30
102/101 [=====] - 162s 2s/step - loss: 0.0036
Epoch 26/30
102/101 [=====] - 162s 2s/step - loss: 0.0035
Epoch 27/30
102/101 [=====] - 164s 2s/step - loss: 0.0035
Epoch 28/30
102/101 [=====] - 163s 2s/step - loss: 0.0034
Epoch 29/30
102/101 [=====] - 162s 2s/step - loss: 0.0033
Epoch 30/30
102/101 [=====] - 163s 2s/step - loss: 0.0035

```

Epoch 00030: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
CPU times: user 1h 36min 6s, sys: 15min 53s, total: 1h 51min 59s
Wall time: 1h 21min 42s

In [11]:

```

model.save(filepath)
model.save_weights("Art_Colorization_Weights.h5")

```

Sample the Results

In [12]:

```

sample = X_test
color_me = gray2rgb(rgb2gray(sample))
color_me_embed = create_inception_embedding(color_me)
color_me = rgb2lab(color_me)[:,:,:0]
color_me = color_me.reshape(color_me.shape+(1,))
output = model.predict([color_me, color_me_embed])
output = output * 128
decoded_imgs = np.zeros((len(output), 256, 256, 3))
for i in range(len(output)):
    cur = np.zeros((256, 256, 3))
    cur[:,:,:0] = color_me[i][:,:,:0]
    cur[:,:,:1:] = output[i]
    decoded_imgs[i] = lab2rgb(cur)
    cv2.imwrite("img_"+str(i)+".jpg", lab2rgb(cur))

```

In [13]:

```

plt.figure(figsize=(20, 6))
for i in range(10):
    # grayscale
    plt.subplot(3, 10, i + 1)
    plt.imshow(rgb2gray(X_test)[i].reshape(256, 256))
    plt.gray()
    plt.axis('off')
    # recolorization
    plt.subplot(3, 10, i + 1 + 10)
    plt.imshow(decoded_imgs[i].reshape(256, 256, 3))
    plt.axis('off')
    # original
    plt.subplot(3, 10, i + 1 + 20)
    plt.imshow(X_test[i].reshape(256, 256, 3))
    plt.axis('off')
plt.tight_layout()
plt.show()

```

